

Comprehensive Project Overview

Real-Time Procedural Resurfacing Using GPU Mesh Shaders

Paper: Raad et al., "Real-time procedural resurfacing using GPU mesh shader," CGF 44(2), Eurographics 2025

Reference implementation: github.com/andarael/resurfacing

1. Project Goal

Reimplement from scratch the GPU mesh shader-based procedural resurfacing framework. The system takes a base mesh M and procedurally generates new geometric surfaces on-the-fly at each face/vertex, using the Vulkan mesh shader pipeline (task shader → mesh shader → fragment shader). No traditional vertex input is used — all geometry is read from storage buffers and generated on-chip.

2. Technology Stack

Component	Choice
Language	C++17 (host), GLSL with GL_EXT_mesh_shader (GPU)
Graphics API	Vulkan 1.3 with VK_EXT_mesh_shader extension
Shader Compilation	glslangValidator → SPIR-V (with GL_GOOGLE_include_directive)
Windowing	GLFW
Math	GLM
UI	ImGui (Vulkan backend)
Mesh Loading	OBJ loader (n-gon support), optionally GLTF for skeletal animation
Build	CMake
Target GPU	NVIDIA RTX 3080 (Ampere), driver 535+
Profiling	NVIDIA Nsight Graphics

The reference implementation uses the Vulkan C++ bindings (`vulkan.hpp`) rather than the C API.

3. High-Level Architecture

3.1 Pipeline Overview (Paper's Framework)

The resurfacing framework has 5 logical stages, mapped to GPU shader stages:



```

F (Mapping):      base mesh -> base elements bi
P (Payload):     culling, LOD, attributes -> ej
K (Amplification): compute mesh shader count kj
Output: taskPayloadSharedEXT + EmitMeshTasksEXT
| (pipelined memory)
v
MESH SHADER (kj invocations per element)
Param: evaluate parametric surface at (u,v)
S:   sample UV grid -> triangulated surface
Output: SetMeshOutputsEXT(verts, prims)
|
v
FRAGMENT SHADER (standard rasterization)
Shading, per-primitive coloring, debug viz

```

3.2 Draw Call

Instead of `vkCmdDraw`, the application calls `vkCmdDrawMeshTasksEXT(cmd, groupCountX, 1, 1)` where `groupCountX = nbFaces + nbVertices` (one task shader workgroup per base element).

3.3 Descriptor Set Layout

Three descriptor sets, derived from the reference implementation:

Set 0 (SceneSet) — per-frame global data:

- Binding 0: ViewUBO (mat4 view, projection, vec4 cameraPosition, float near, far)
- Binding 1: GlobalShadingUBO (light position, ambient, shading params)

Set 1 (HESet) — half-edge mesh data (Structure of Arrays):

- Binding 0: vec4 buffers array[5] (vertex positions, colors, normals, face normals, face centers)
- Binding 1: vec2 buffers array[1] (vertex texcoords)
- Binding 2: int buffers array[10] (vertex edges, face edges, face vert counts, face offsets, HE vertex/face/next/prev/twin, vertex-face indices)
- Binding 3: float buffers array[1] (face areas)

Set 2 (PerObjectSet) — per-object data:

- Binding 0: Config UBO (ResurfacingUBO or PebbleUBO)
- Binding 1: ShadingUBO (per-object shading)
- Binding 2: LUT vertex buffer (control cage points for B-spline)
- Bindings 3–5: Skeletal animation buffers (joint indices, weights, bone matrices)
- Binding 6: Samplers[2] (linear, nearest)
- Binding 7: Textures[2] (AO texture, element type texture/control map)

Push constants: model matrix (mat4), MVP (mat4).

4. Mesh Data Structure

4.1 Input Format (N-gon OBJ)

The system loads OBJ files supporting n-gon faces (not just triangles). Each face stores:

```
struct NGonFace {
    vec4 normal;      // face normal
    vec4 center;     // face centroid
    uint32 offset;   // offset into index array
    uint32 count;    // vertex count (3, 4, 5, ...)
    float faceArea; // area of the face
};
```

4.2 Half-Edge Conversion

The loaded ngon mesh is converted to a half-edge data structure on the CPU using Structure of Arrays (SoA) layout:

```
struct HalfEdgeMesh {
    uint32 nbVertices, nbFaces;
    // Vertex SoA:
    vector<vec4> positions, colors, normals;
    vector<vec2> texCoords;
    vector<int> edges;           // one outgoing half-edge per vertex
    // Face SoA:
    vector<int> faceEdges, vertCounts, offsets;
    vector<vec4> faceNormals, faceCenters;
    vector<float> faceAreas;
    // Half-edge SoA:
    vector<int> heVertex, heFace, heNext, hePrev, heTwin;
    vector<int> vertexFaceIndices; // flattened per-face vertex lists
};
```

4.3 GPU Upload and Shader Access

Each array is uploaded as a separate SSBO. The shader accesses them via typed buffer arrays:

```
layout(std430, set=1, binding=0) readonly buffer { vec4 data[]; } heVec4Buffer[5];
layout(std430, set=1, binding=2) readonly buffer { int data[]; } heIntBuffer[10];

vec3 getVertexPosition(uint vertId) { return heVec4Buffer[0].data[vertId].xyz; }
vec3 getFaceNormal(uint faceId)    { return heVec4Buffer[3].data[faceId].xyz; }
uint getHalfEdgeNext(uint edgeId)   { return heIntBuffer[6].data[edgeId]; }
```

5. Task Shader (F + P + K)

5.1 Mapping Function F

One task shader workgroup is dispatched per base element. The workgroup ID determines whether it's a face or vertex element:

```

uint faceId = gl_WorkGroupID.x;
bool isVertex = (faceId >= nbFaces);
if (isVertex) {
    vertId = faceId - nbFaces;
    faceId = getFaceId(vertId); // get face via half-edge
}

```

5.2 Payload Function P

Culling: Back-face culling via $\text{dot}(\text{viewDir}, \text{instanceNormal}) > \text{threshold}$. Frustum culling by projecting center to clip space and checking bounds.

Skinning (optional): If skeletal animation data is bound, apply bone transforms to position/normal before culling.

Element type selection: If a control map texture is bound, sample it to determine element type per face. Otherwise use global elementType from UBO.

5.3 LOD Computation

```

uvec2 getLodMN(LodInfos info, uint elementType) {
    if (!doLod) return MN;
    parametricBoundingBox(info, elementType);
    float screenSpaceSize = boundingBoxScreenSpaceSize(info);
    uvec2 targetMN = uvec2(MN * sqrt(screenSpaceSize * lodFactor));
    return clamp(targetMN, minResolution, MN);
}

```

5.4 Amplification Function K

Due to hardware limits (`MAX_VERTICES`=81, `MAX_PRIMITIVES`=128 for 8x8 grids), the UV domain is split across multiple mesh shader invocations:

```

uvec2 getDeltaUV(uvec2 MN) {
    // Find largest deltaU x deltaV fitting hardware limits
    // (deltaU+1)*(deltaV+1) <= MAX_VERTICES
    // deltaU*deltaV*2 <= MAX_PRIMITIVES
    for (deltaU = max..1) for (deltaV = max..1)
        if (fits) return uvec2(deltaU, deltaV);
}
uvec2 numMeshTasks = (MN + deltaUV - 1) / deltaUV;

```

5.5 Task Payload Structure

```

struct TaskPayload {
    vec3 position; // world-space center
    vec3 normal; // orientation normal
    float area; // face area (for scaling)
    uint taskId; // original workgroup ID
    bool isVertex; // face vs vertex element
    uvec2 MN; // target resolution
    uvec2 deltaUV; // per-invocation UV range
    uint elementType; // which parametric surface
};

```

```
taskPayloadSharedEXT TaskPayload taskPayload;
```

5.6 Emit

```
EmitMeshTasksEXT(numMeshTasks.x * doRender, numMeshTasks.y * doRender, 1);
```

Setting count to 0 skips the element (culled).

6. Mesh Shader (Param + S)

6.1 Configuration

```
layout(local_size_x = 32) in; // MESH_GROUP_SIZE = 32
layout(max_vertices = 81, max_primitives = 128, triangles) out;
```

6.2 UV Domain Subsection

```
uvec2 startUV = uvec2(gl_WorkGroupID.xy * deltaUV.xy);
uvec2 localDeltaUV = min(deltaUV, MN - startUV);
uint numVertices = (localDeltaUV.x + 1) * (localDeltaUV.y + 1);
uint numPrimitives = localDeltaUV.x * localDeltaUV.y * 2;
SetMeshOutputsEXT(numVertices, numPrimitives);
```

6.3 Surface Sampling (S)

```
for (uint u = gl_LocalInvocationID.x; u <= localDeltaUV.x; u += MESH_GROUP_SIZE) {
    for (uint v = 0; v <= localDeltaUV.y; ++v) {
        vec2 uvCoords = vec2(startUV + uvec2(u, v)) / vec2(MN);
        vec3 pos, normal;
        parametricPosition(uvCoords, pos, normal, elementType);
        offsetVertex(pos, normal); // scale, rotate, translate
        emitVertex(pos, normal, uvCoords, vertexIndex);
    }
}
```

6.4 Triangle Emission

```
gl_PrimitiveTriangleIndicesEXT[2*q+0] = uvec3(v00, v10, v11);
gl_PrimitiveTriangleIndicesEXT[2*q+1] = uvec3(v00, v11, v01);
```

6.5 Vertex Transform (offsetVertex)

1. **Scale** by $\text{sqrt}(\text{faceArea}) * \text{userScaling}$
2. **Rotate** to align with face/vertex normal using `align_rotation_to_vector()`
3. **Translate** to element position (face center or vertex position)

7. Parametric Surfaces

7.1 Analytical Surfaces

Implemented via switch on `elementType` (0–8):

- Type 0: Torus — `parametricTorus(uv, pos, nrm, majorR, minorR)`

- Type 1: Sphere — parametricSphere(uv, pos, nrm, radius)
- Type 2: Möbius Strip
- Type 3: Klein Bottle
- Type 4: Hyperbolic Paraboloid
- Type 5: Helicoid
- Type 6: Cone
- Type 7: Cylinder
- Type 8: Egg

Each function also provides a bounding box function for LOD screen-space size estimation.

7.2 B-Spline Control Cages (types 9–10)

Control cage vertices loaded from separate OBJ (quad grid, sorted by UV), uploaded as SSBO. The mesh shader evaluates bicubic B-spline or Bezier surfaces with configurable degree and cyclic boundary conditions.

```
struct LutData {
    vector<vec4> positions; // control points flattened row-major
    uint Nx, Ny;           // grid dimensions
    vec3 min, max;         // bounding box of control cage
};
```

7.3 Pebbles (Procedural Cages)

Pebbles use a completely separate pipeline with different payload and generation logic:

Task shader: Maps one element per face only. Fetches face vertex positions into shared memory. Emits $\text{vertCount} \times 2 \times 3$ patches per face (edge regions \times extrusion layers), plus fan/ring patches.

Mesh shader: Constructs control cage on-the-fly: extrude vertices along face normal, project toward center for roundness, subdivide once to get 16 B-spline control points in shared memory. Evaluates bicubic B-spline per patch. Applies procedural noise (Gabor/Perlin).

Pebble UBO parameters: subdivisionLevel (0–8), extrusionAmount, extrusionVariation, roundness (0–2), doNoise, noiseAmplitude, noiseFrequency.

8. Control Maps (Hybrid Surfaces)

Per-face element type selection via a color-coded texture sampled in the task shader using base mesh UV coordinates. Hard-coded color mapping (blue \rightarrow spike, violet \rightarrow ball, green \rightarrow empty). Elements with type=-1 are skipped (base mesh shows through).

9. Shading

9.1 Fragment Shader

Standard Blinn-Phong shading with per-primitive ID coloring, AO texture support, configurable ambient/diffuse/specular/shininess, and HSV-based random coloring from element ID.

9.2 Mesh-to-Fragment Data

- Per-vertex: worldPosU (vec4: xyz=world pos, w=u), normalV (vec4: xyz=normal, w=v)
- Per-primitive: data (uvec4: task ID, face area, debug)

10. Vulkan Renderer Infrastructure

10.1 Pipeline Creation

- Auto-detect shader stage from file extension (.task, .mesh, .frag)
- pVertexInputState = nullptr, pInputAssemblyState = nullptr
- Dynamic rendering (VK_KHR_dynamic_rendering) instead of render passes
- Shader compilation: glslangValidator --target-env vulkan1.3

10.2 Required Extensions and Features

```
// Device extensions
VK_KHR_SWAPCHAIN_EXTENSION_NAME
VK_EXT_MESH_SHADER_EXTENSION_NAME

// Device features
VkPhysicalDeviceMeshShaderFeaturesEXT { .taskShader=TRUE, .meshShader=TRUE }
```

10.3 Frame Loop

```
beginFrame() -> beginRendering() ->
    bind parametric pipeline -> for each object: bind + dispatch
        bind HE pipeline -> for each base mesh: bind + dispatch
            bind pebble pipeline -> for each pebble: bind + dispatch
    -> endRendering() -> renderUI() -> endFrame()
```

Each object's bindAndDispatch: bind HE descriptors (set 1), bind per-object descriptors (set 2), push constants (model matrix), call vkCmdDrawMeshTasksEXT(cmd, nbFaces + nbVertices, 1, 1).

11. Implementation Tasks (Ordered)

Phase 1: Vulkan Infrastructure (Weeks 1–2)

Task 1.1: Project Scaffold

- CMakeLists.txt with Vulkan, GLFW, GLM, Dear ImGui dependencies
- Shader compilation rules (glslangValidator with include directive support)
- Window creation, Vulkan instance (API 1.3), surface

Task 1.2: Device Setup

- Physical device selection with VK_EXT_mesh_shader support
- Logical device creation with mesh shader features enabled via pNext chain
- Queue families, command pool
- Print mesh shader hardware limits (maxMeshOutputVertices, maxMeshOutputPrimitives, maxTaskPayloadSize)

Task 1.3: Swap Chain and Rendering

- Swap chain creation, depth buffer (D32_SFLOAT)
- Dynamic rendering or render pass setup, frame synchronization
- Basic clear-screen render loop

Task 1.4: Descriptor Infrastructure

- Create 3 descriptor set layouts matching Section 3.3
- Descriptor pool, allocate descriptor sets
- Create and map uniform buffers for ViewUBO and GlobalShadingUBO

Task 1.5: Mesh Shader Pipeline Object

- Graphics pipeline with task + mesh + fragment stages, no vertex input
- Load vkCmdDrawMeshTasksEXT via vkGetDeviceProcAddr
- Test with hardcoded triangle output from mesh shader

Task 1.6: ImGui Integration

- ImGui Vulkan backend setup, basic UI with camera controls and FPS

Phase 2: Mesh Loading and GPU Upload (Weeks 2–3)

Task 2.1: OBJ Loader with N-gon Support

- Parse OBJ with arbitrary face vertex counts, compute per-face normals/centers/areas

Task 2.2: Half-Edge Construction

- Convert NgonData → HalfEdgeMesh (SoA layout), build twin map, compute vertexFaceIndices

Task 2.3: GPU Buffer Upload

- Upload each SoA array as separate SSBO, update HE descriptor set bindings

Task 2.4: Shared Shader Interface

- Create shaderInterface.h shared between C++ and GLSL with #ifdef __cplusplus guards

Phase 3: Core Resurfacing Pipeline (Weeks 3–4)

Task 3.1: Task Shader — Mapping Function F

- Dispatch per face+vertex, read position/normal from SSBO, populate payload, emit 1 mesh invocation

Task 3.2: Mesh Shader — Hardcoded Geometry

- Output single quad at payload position to validate task→mesh→fragment data path

Task 3.3: Parametric Surface Evaluation

- Implement parametricTorus in GLSL, evaluate 8x8 UV grid, apply scale/rotate/translate

Task 3.4: Multiple Parametric Types

- Add sphere, cone, cylinder, egg; ImGui control for type and radii

Task 3.5: Proper UV Grid Emission

- Grid vertices with quad→triangle pairs, per-vertex normals/UVs, per-primitive IDs

Phase 4: Amplification and LOD (Weeks 4–5)

Task 4.1: Amplification Function K

- getDeltaUV() for hardware-constrained UV tiles, 2D EmitMeshTasksEXT

Task 4.2: Variable Resolution

- Per-axis MN via UBO, handle edge tiles in mesh shader

Task 4.3: Pipeline Permutations (Optional)

- Multiple compiled variants (4x4 small, 8x8 large) for optimal GPU utilization

Task 4.4: Frustum Culling

- Project center to clip space, check bounds, set count=0 for culled

Task 4.5: Back-Face Culling

- Normal–view direction dot product with configurable threshold

Task 4.6: Screen-Space LOD

- Parametric bounding box → screen space projection → adaptive MN resolution

Phase 5: B-Spline Control Cages (Weeks 5–7)

Task 5.1: LUT Loader

- Load control cage OBJ (quad grid), sort by UV, flatten, upload as SSBO

Task 5.2: B-Spline Evaluation

- Cubic B-spline basis functions, 4x4 control point neighborhoods, cyclic boundaries

Task 5.3: Bezier Evaluation

- Bernstein basis with configurable degree (1–3)

Task 5.4: Control Cage LOD

- Bounding box from LUT extents, same screen-space LOD framework

Phase 6: Pebble Generation (Weeks 7–9)

Task 6.1: Pebble Pipeline

- Separate pipeline + PebbleUBO, one task per face only

Task 6.2: Pebble Task Shader

- Fetch face vertices to shared memory, compute patch count and subdivision level

Task 6.3: Pebble Control Cage Construction

- Extrude, project for roundness, subdivide once → 16 control points in shared memory

Task 6.4: Pebble Surface Evaluation

- Bicubic B-spline from procedural cage, procedural noise, random variation

Task 6.5: Pebble LOD

- Bounding box from extruded vertices, power-of-two subdivision levels

Phase 7: Control Maps and Polish (Weeks 9–10)

Task 7.1: Control Map Textures

- Sample color-coded texture in task shader, map to element types

Task 7.2: Base Mesh Rendering

- Simple mesh shader pipeline for wireframe/solid base mesh visualization

Task 7.3: Skeletal Animation (Optional)

- GLTF loader, skin transform in task shader before culling

Phase 8: Performance Analysis (Weeks 10–12)

Task 8.1: GPU Timing

- Vulkan timestamp queries, per-pipeline timing, ImGui display

Task 8.2: Benchmarking Suite

- Reproduce Table 1: frame times per surface type, with/without culling/LOD

Task 8.3: Memory Analysis

- VRAM usage comparison: mesh shader vs static vertex buffers

Task 8.4: Comparison Pipelines (Optional)

- Vertex pipeline and tessellation pipeline baselines for Figure 7 reproduction

12. Key Implementation Gotchas

- 1.** pVertexInputState and pInputAssemblyState MUST be nullptr in the mesh shader pipeline. Setting them causes validation errors or crashes.
- 2.** taskPayloadSharedEXT is the ONLY way to pass data from task → mesh. Payload size limited to ~16KB. Keep it lean.
- 3.** Shader include directives require GL_GOOGLE_include_directive. Pass -l flag to glslangValidator.
- 4.** gl_WorkGroupID in mesh shader: when EmitMeshTasksEXT emits (x, y, 1), mesh shader gets gl_WorkGroupID.x and .y for UV tile indices.
- 5.** MAX_VERTICES and MAX_PRIMITIVES are compile-time constants. Cannot change at runtime. Use pipeline permutation strategy.
- 6.** SoA half-edge layout uses arrays of typed buffers (heVec4Buffer[5], heIntBuffer[10]). Requires GL_EXT_scalar_block_layout and matching descriptor array counts.
- 7.** Pebble mesh shader: control cage too large for payload. Must re-evaluate from SSBO in mesh shader, store in shared memory.
- 8.** Push constants for model matrix (changes per draw call), not UBO.
- 9.** Vulkan Y-flip: projection matrix needs proj[1][1] *= -1.
- 10.** Subgroup operations (GL_KHR_shader_subgroup_arithmetic) enabled but used minimally. Useful for cooperative vertex emission.

13. File Structure (Reference Implementation)

```

resurfacing/
+-- CMakeLists.txt
+-- assets/
|   +-- demo/                      # Dragon model, textures, animations
|   +-- icosphere.obj              # Test mesh
|   +-- parametric_luts/           # Control cage OBJ files
+-- libs/                          # glfw, glm, imgui, stb, tinygltf
+-- shaders/
|   +-- shaderInterface.h          # Shared CPU/GPU: bindings, UBOs, getters
|   +-- common.glsl                 # lerp, isVisible, rotation alignment
|   +-- lods.glsl                  # LOD: bounding box -> screen space
|   +-- noise.glsl                 # Procedural noise (Gabor, Perlin)
|   +-- shading.glsl                # Fragment shading
|   +-- stdPerVertexMesh.glsl       # Per-vertex/per-primitive output interface
|   +-- parametric/
|       +-- parametric.glsl         # Payload, constants, LOD, amplification
|       +-- parametricSurfaces.glsl # Analytical surface equations
|       +-- parametricGrids.glsl    # B-spline and Bezier evaluation
|       +-- parametric.task/.mesh/.frag
+-- pebbles/
|   +-- pebble.glsl/.task/.mesh/.frag
+-- halfEdges/
|   +-- halfEdge.mesh, halfedge.frag
+-- src/
|   +-- main.cpp                   # App, init, render loop, UI
|   +-- renderer.cpp/.hpp          # Vulkan wrapper
|   +-- vkHelper.hpp               # Buffer, Texture, Pipeline types
|   +-- HalfEdge.hpp               # Half-edge structure + conversion
|   +-- AppRessources.cpp/.hpp     # Scene objects, buffer management
|   +-- loaders/ObjLoader, GLTFLoader

```

14. Success Criteria

4. **Tori/spheres/cones** render correctly on loaded OBJ meshes (chain mail appearance)
5. **B-spline control cages** produce smooth subdivision surfaces (scale appearance)
6. **Pebbles** generate from face geometry with noise perturbation
7. **Culling** eliminates off-screen and back-facing elements
8. **LOD** varies resolution based on screen-space size
9. **Control maps** enable per-face element type selection
10. **Performance** within same order of magnitude as paper's results on RTX 3080
11. **Memory footprint** is constant regardless of subdivision level (geometry never stored)